

# CS 556 – Distributed Systems

## Tutorial on Java Threads

4 Oct 2002

Java Threads

# Threads

- A thread is a lightweight process – a single sequential flow of execution within a program
- Threads make possible the implementation of programs that seem to perform multiple tasks at the same time (e.g. multi-threaded Web servers)
- A new way to think about programming

# Java Threads

We will cover:

- How to create threads in Java
- The Life Cycle of a thread
- Thread Priority
- Synchronization of threads
- Grouping of threads

# How to create Java Threads

There are two ways to create a Java thread:

1. Extend the `java.lang.Thread` class
2. Implement the `java.lang.Runnable` interface

# Extending the Thread class

- In order to create a new thread we may subclass `java.lang.Thread` and customize what the thread does by overriding its empty `run` method.
- The `run` method is where the action of the thread takes place.
- The execution of a thread starts by calling the `start` method.

# Example I

```
class MyThread extends Thread {  
    private String name, msg;  
  
    public MyThread(String name, String msg) {  
        this.name = name;  
        this.msg = msg;  
    }  
    public void run() {  
        System.out.println(name + " starts its execution");  
        for (int i = 0; i < 5; i++) {  
            System.out.println(name + " says: " + msg);  
            try {  
                Thread.sleep(5000);  
            } catch (InterruptedException ie) {}  
        }  
        System.out.println(name + " finished execution");  
    }  
}
```

# Example I

```
class MyThread extends Thread {  
    private String name, msg;  
  
    public MyThread(String name, String msg) {  
        this.name = name;  
        this.msg = msg;  
    }  
    public void run() {  
        System.out.println(name + " starts its execution");  
        for (int i = 0; i < 5; i++) {  
            System.out.println(name + " says: " + msg);  
            try {  
                Thread.sleep(5000);  
            } catch (InterruptedException ie) {}  
        }  
        System.out.println(name + " finished execution");  
    }  
}
```

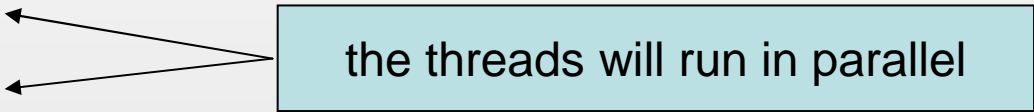
# Example I

```
class MyThread extends Thread {  
    private String name, msg;  
  
    public MyThread(String name, String msg) {  
        this.name = name;  
        this.msg = msg;  
    }  
    public void run() {  
        System.out.println(name + " starts its execution");  
        for (int i = 0; i < 5; i++) {  
            System.out.println(name + " says: " + msg);  
            try {  
                Thread.sleep(5000);  
            } catch (InterruptedException ie) {}  
        }  
        System.out.println(name + " finished execution");  
    }  
}
```



# Example I (cont.)

```
public class test {  
    public static void main(String[] args) {  
        MyThread mt1 = new MyThread("thread1", "ping");  
        MyThread mt2 = new MyThread("thread2", "pong");  
        mt1.start();  
        mt2.start();  
    }  
}
```



the threads will run in parallel

# Example I (cont.)

- Typical output of the previous example:

thread1 starts its execution

thread1 says: ping

thread2 starts its execution

thread2 says: pong

thread1 says: ping

thread2 says: pong

thread1 says: ping

thread2 says: pong

thread1 says: ping

thread2 says: pong

thread1 says: ping

thread2 says: pong

thread1 finished execution

thread2 finished execution

# Implementing the Runnable interface

- In order to create a new thread we may also provide a class that implements the `java.lang.Runnable` interface
- Preferred way in case our class has to subclass some other class
- A Runnable object can be wrapped up into a Thread object
  - `Thread(Runnable target)`
  - `Thread(Runnable target, String name)`
- The thread's logic is included inside the run method of the runnable object

# Example II

```
class MyClass implements Runnable {  
    private String name;  
    private A sharedObj;  
    public MyClass(String name, A sharedObj) {  
        this.name = name; this.sharedObj = sharedObj;  
    }  
    public void run() {  
        System.out.println(name + " starts execution");  
        for (int i = 0; i < 5; i++) {  
            System.out.println(name + " says: " + sharedObj.getValue());  
            try {  
                Thread.sleep(5000);  
            } catch (InterruptedException ie) {}  
        }  
        System.out.println(name + " finished execution");  
    }  
}
```

## Example II (cont.)

```
class A {  
    private String value;  
    public A(String value) { this.value = value; }  
    public String getValue() {  
        return value;  
    }  
}
```

```
public class test2 {  
    public static void main(String[] args) {  
        A sharedObj = new A("some value");  
        Thread mt1 = new Thread(new MyClass("thread1", sharedObj));  
        Thread mt2 = new Thread(new MyClass("thread2", sharedObj));  
        mt1.start(); mt2.start();  
    }  
}
```

shared variable



# Example II (cont.)

- Typical output of the previous example:

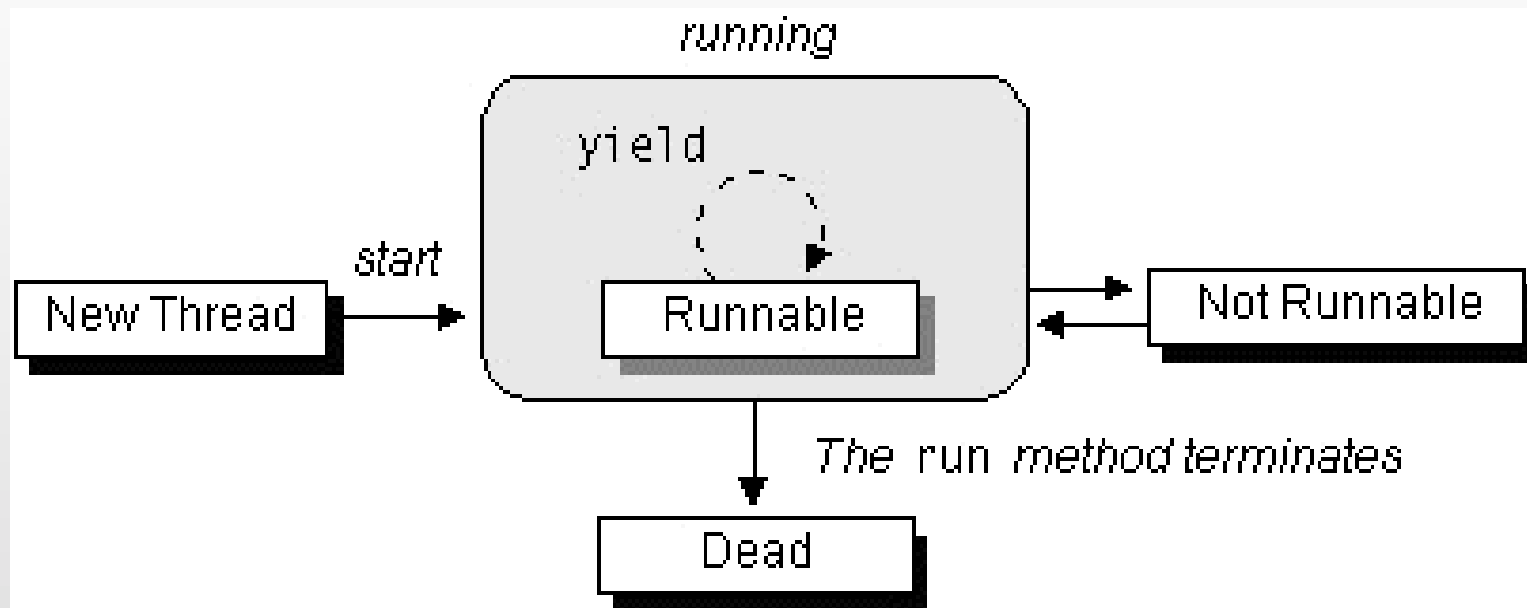
```
thread1 starts execution
thread1 says: some value
thread2 starts execution
thread2 says: some value
thread1 says: some value
thread2 says: some value
thread1 says: some value
thread2 says: some value
thread1 says: some value
thread2 says: some value
thread1 says: some value
thread2 says: some value
thread1 finished execution
thread2 finished execution
```

# Java Threads

We will cover:

- How to create threads in Java
- **The Life Cycle of a thread**
- Thread Priority
- Synchronization of threads
- Grouping of threads

# The Life Cycle of a Thread





# The Life Cycle of a Thread (cont.)

- The start method creates the system resources necessary to run the thread, schedules the thread to run, and calls the thread's run method
- A thread becomes Not Runnable when one of these events occurs:
  - Its `sleep` method is invoked.
  - The thread calls the `wait` method.
  - The thread is blocking on I/O
- A thread dies naturally when the `run` method exits

# Java Threads

We will cover:

- How to create threads in Java
- The Life Cycle of a thread
- **Thread Priority**
- Synchronization of threads
- Grouping of threads

# Thread Priority

- On a single CPU threads actually run one at a time in such a way as to provide an illusion of concurrency.
- Execution of multiple threads on a single CPU, in some order, is called *scheduling*.
- The Java runtime supports a very simple scheduling algorithm (fixed priority scheduling). This algorithm schedules threads based on their priority relative to other runnable threads.

# Thread Priority (cont.)

- The runtime system chooses the runnable thread with the highest priority for execution
- If two threads of the same priority are waiting for the CPU, the scheduler chooses one of them to run in a round-robin fashion.
- The chosen thread will run until:
  - A higher priority thread becomes runnable.
  - It yields, or its `run` method exits.
  - On systems that support time-slicing, its time allotment has expired

# Thread Priority (cont.)

- When a Java thread is created, it inherits its priority from the thread that created it.
- You can modify a thread's priority at any time after its creation using the `setPriority` method.
- Thread priorities are integers ranging between `MIN_PRIORITY` and `MAX_PRIORITY` (constants defined in the `Thread` class). The higher the integer, the higher the priority.

# Thread Priority (cont.)

- Use priority only to affect scheduling policy for efficiency purposes.
- Do **not** rely on thread priority for algorithm correctness.

# Java Threads

We will cover:

- How to create threads in Java
- The Life Cycle of a thread
- Thread Priority
- **Synchronization of threads**
- Grouping of threads

# Synchronization of Threads

- In many cases concurrently running threads share data and must consider the state and activities of other threads.
- If two threads can both execute a method that modifies the state of an object then the method should be declared to be `synchronized`, allowing only one thread to execute the method at a time.
- If a class has at least one synchronized methods, each instance of it has a monitor. A monitor is an object that can block threads and notify them when it is available.



# Synchronization of Threads (cont.)

Example:

```
public synchronized void updateRecord() {  
    // critical code goes here ...  
}
```

- Only one thread may be inside the body of this function. A second call will be blocked until the first call returns or `wait()` is called inside the synchronized method.

# Synchronization of Threads (cont.)

- If you don't need to protect an entire method, you can synchronize on an object:

```
public void foo() {  
    //...  
    synchronized (this) {  
        //critical code goes here ...  
    }  
    //...  
}
```

- Declaring a method as `synchronized` is equivalent to synchronizing on *this* for all the method block.

# Synchronization of Threads (cont.)

- The Object class has three synchronization methods:
  - wait()
  - notify()
  - notifyAll()
- These methods allow objects to wait until another object notifies them:

```
synchronized( waitForThis ) {  
    try {  
        waitForThis.wait();  
    } catch (InterruptedException ie) {}  
}
```

- To wait on an object, you must first synchronize on it.

# Synchronization of Threads (cont.)

- A thread may call `wait()` inside a synchronized method. A timeout may be provided. If missing or zero then the thread waits until either `notify()` or `notifyAll()` is called, otherwise until the timeout period expires.
- `wait()` is called by the thread owning the lock associated with a particular object (it causes the lock to be released).
- `notify()` or `notifyAll()` are only called from a synchronized method. One or all waiting threads are notified, respectively. It's probably better/safer to use `notifyAll()`. These methods don't release the lock. The threads awakened will not return from their `wait()` call immediately, but only when the thread that called `notify()` or `notifyAll()` finally relinquishes ownership of the lock.

# Synchronization of Threads (cont.)

## A Producer – Consumer Problem

```
public synchronized void put(int value)
{
    while (available == true) {
        try {
            //wait for Consumer to get
            //value
            wait();
        } catch (InterruptedException e)
        {}
    }
    contents = value;
    available = true;
    //notify Consumer that value has
    //been set
    notifyAll();
}

public synchronized int get() {
    while (available == false) {
        try {
            //wait for Producer to put
            //value
            wait();
        } catch (InterruptedException e)
        {}
    }
    available = false;
    //notify Producer that value has
    //been retrieved
    notifyAll();
    return contents;
}
```

# Synchronization of Threads (cont.)

- The Thread class has a method, `join()`, which allows an object to wait until the thread terminates

```
public void myMethod() {  
    // Do some work here...  
    // Can't proceed until another thread is done:  
    otherThread.join();  
    // Continue work...  
}
```

- Equivalent to `waitFor()` of `java.lang.Process`

# Java Threads

We will cover:

- How to create threads in Java
- The Life Cycle of a thread
- Thread Priority
- Synchronization of threads
- **Grouping of threads**

# Grouping of Threads

- Thread groups provide a mechanism for collecting multiple threads into a single object and manipulating those threads all at once, rather than individually.
- Java thread groups are implemented by the `java.lang.ThreadGroup` class
- When a Java application first starts up, the Java runtime system creates a `ThreadGroup` named `main`.
- Unless specified otherwise, all new threads that you create become members of the main thread group.



# Grouping of Threads (cont.)

- To put a new thread in a thread group the group must be explicitly specified when the thread is created
  - public **Thread**(ThreadGroup *group*, Runnable *runnable*)
  - public **Thread**(ThreadGroup *group*, String *name*)
  - public **Thread**(ThreadGroup *group*, Runnable *runnable*, String *name*)
- A thread can not be moved to a new group after the thread has been created.

# Grouping of Threads (cont.)

- A ThreadGroup may also contain other ThreadGroups allowing the creation of an hierarchy of threads and thread groups:
  - `public ThreadGroup(ThreadGroup parent, String name)`
- To get the thread group of a thread the `getThreadGroup` of the Thread class should be called:

*ThreadGroup group = myThread.getThreadGroup();*

# Reference

- Read the Java Tutorial on Threads:

<http://java.sun.com/docs/books/tutorial/essential/threads/>